

A Hitchhiker's Guide to Reinventing a Prolog Machine

Paul Tarau

Department of Computer Science and Engineering
University of North Texas

ICLP'17

Deriving the execution algorithm

Prolog: the two-clause meta-interpreter

The meta-interpreter `metaint/1` uses a (difference)-list view of prolog clauses.

```
metaint([]).           % no more goals left, succeed
metaint([G|Gs]):-      % unify the first goal with the head of a clause
    cls([G|Bs],Gs),    % build a new list of goals from the body of the
                      % clause extended with the remaining goals as tail
    metaint(Bs).       % interpret the extended body
```

- clauses are represented as facts of the form `cls/2`
- the first argument representing the head of the clause + a list of body goals
- clauses are terminated with a variable, also the second argument of `cls/2`.

```
cls([ add(0,X,X)                               |Tail],Tail).
cls([ add(s(X),Y,s(Z)), add(X,Y,Z)             |Tail],Tail).
cls([ goal(R), add(s(s(0)),s(s(0)),R)          |Tail],Tail).
```

```
?- metaint([goal(R)]).
R = s(s(s(s(0)))) .
```

The “natural language equivalent” of the equational form

As the recursive tree structure of a Prolog is flattened, it makes sense to express it as an equivalent “natural language-looking” sentence.

```
add SX Y SZ if SX holds s X and SZ holds s Z and add X Y Z.
```

- note and the correspondence between the keywords “if” and “and” to Prolog’s “:-” clause neck and “,” conjunction symbols
- the keyword “holds” (like the use of Prolog’s “=”) expresses a unification operation between a variable and a flattened Prolog term
- the toplevel skeleton of the clause can be kept implicit as it is easily recoverable
- this is our “assembler language” to be read in directly by the loader of a runtime system
- a simple tokenizer splitting into words sentences delimited by “.” is all that is needed to complete a parser for this English-style “assembler language”

The heap representation as executable code

Representing terms in a Java-based runtime

- we instruct our tokenizer to recognize variables, symbols and (small) integers as primitive data types
- we develop a Java-based interpreter in which we represent our Prolog terms top-down
- Java's primitive `int` type is used for tagged words
- in a **C** implementation one might want to chose `long long` instead of `int` to take advantage of the 64 bit address space
- we instruct our parser to extract as much information as possible by marking each word with a relevant tag

The top-down representation of terms

`add(s(X), Y, s(Z)) :- add(X, Y, Z) .`

compiles to

`add _0 Y _1 and _0 holds s X and _1 holds s Z if add X Y Z .`

- on the heap (starting in this case at address 5):

```
[5]a:4 [6]c:add [7]r:10 [8]v:8 [9]r:13 [10]a:2 [11]c:s [12]v:12
[13]a:2 [14]c:s [15]v:15 [16]a:4 [17]c:add [18]u:12 [19]u:8
[20]u:15
```

- distinct tags of first occurrences (tagged “v:”) and subsequent occurrences of variables (tagged “u:”).
- references (tagged “r:”) always point to arrays starting with their length marked with tag “a:”
- cells tagged as array length contain the arity of the corresponding function symbol incremented by 1
- the “skeleton” of the clause in the previous example is shown as:

```
r:5 :- [r:16]
```

Clauses as descriptors of heap cells

- the parser places the cells composing a clause directly to the heap
- a descriptor (defined by the small class `Clause`) is created and collected to the array called “`clauses`” by the parser
- an object of type `Clause` contains the following fields:
 - `int base`: the base of the heap where the cells for the clause start
 - `int len`: the length of the code of the clause i.e., number of the heap cells the clause occupies
 - `int neck`: the length of the head and thus the offset where the first body element starts (or the end of the clause if none)
 - `int[] gs`: the toplevel skeleton of a clause containing references to the location of its head and then body elements
 - `int[] xs`: the index vector containing dereferenced constants, numbers or array sizes as extracted from the outermost term of the head of the clause, with 0 values marking variable positions.

Execution as iterated clause unfolding

The key intuition: we emulate (procedurally) the meta-interpreter

- as the meta-interpreter shows it, Prolog's execution algorithm can be seen as iterated unfolding of a goal with heads of matching clauses
- if unification is successful, we extend the list of goals with the elements of the body of the clause, to be solved first
- as we do not assume anymore that predicate symbols are non-variables, it makes sense to design indexing as a distinct
- pre-unification step: detecting matching clauses without copying to the heap
- one can filter matching clauses by comparing the outermost array of the current goal with the outermost array of a clause head
- we use for this the prototype of a clause head without starting to place new terms on the heap
- dereferencing is avoided when working with material from the heap-represented clauses (none for first occurrences of variables, exactly once for others)

Fast “linear” term relocation and the immutable goal stack

- we implement a fast relocation loop that speculatively places the clause head (including its subterms) on the heap
- this “single instruction multiple data” operation can benefit from parallel execution!
- new terms are built on the heap by the relocation loop in two stages: first the clause head (including its subterms) and then, if unification succeeds, also the body
- stretching out the Spine: the (immutable) goal stack
 - a `Spine` is a runtime abstraction of a `Clause`
 - it collects information needed for the execution of the goals originating from it
 - goal elements on this immutable list are shared among alternative branches

The execution algorithm

Our interpreter: yielding an answer and ready to resume

- it starts from a `Spine` and works through a stream of answers, returned to the caller one at a time, until the `spines` stack is empty
- it returns null when no more answers are available

```
final Spine yield() {
    while (!spines.isEmpty()) {
        final Spine G = spines.peek();
        if (hasClauses(G)) {
            if (hasGoals(G)) {
                final Spine C = unfold(G);
                if (C != null) {
                    if (!hasGoals(C)) return C; // return answer
                    else spines.push(C);
                } else popSpine(); // no matches
            } else unwindTrail(G.ttop); // no more goals in G
        } else popSpine(); // no clauses left
    }
    return null;
}
```

Exposing the answers of a logic engine to the implementation language

Answer streams

- to encapsulate our answer streams in a Java 8 `stream`, a special iterator-like interface called `Splititerator` is used
- the work is done by the `tryAdvance` method which yields answers while they are not equal to `null`, and terminates the `stream` otherwise

```
public boolean tryAdvance(Consumer<Object> action) {  
    Object R = ask();  
    boolean ok = null != R;  
    if (ok) action.accept(R);  
    return ok;  
}
```

- three more methods are required by the interface, mostly to specify when to stop the stream and that the stream is ordered and sequential

Multi-argument indexing: a modular add-on

The indexing algorithm

- the indexing algorithm is designed as an independent add-on to be plugged into the the main Prolog engine
- for each argument position in the head of a clause it associates to each indexable element (symbol, number or arity) the set of clauses where the indexable element occurs in that argument position
- the clauses having variables in an indexed argument position are also collected in a separate set for each argument position
- 3 levels are used, closely following the data that we want to index
- sets of clause numbers associated to each (tagged) indexable element are backed by an `IntMap` implemented as a fast `int-to-int` hash table (using linear probing)
- an `IntMap` is associated to each indexable element by a `HashMap`
- the `HashMap`s are placed into an array indexed by the argument position to which they apply

– continued –

- when looking for the clauses matching an element of the list of goals to solve, for an indexing element x occurring in position i , we fetch the the set $C_{x,i}$ of clauses associated to it
- If V_i denotes the set of clauses having variables in position i , then any of them can also unify with our goal element
- thus we would need to compute the union of the sets $C_{x,i}$ and V_i for each position i , and then intersect them to obtain the set of matching clauses
- instead of actually compute the unions for each element of the set of clauses corresponding to the “predicate name” (position 0), we retain only those which are either in $C_{x,i}$ or in V_i for each $i > 0$
- we do the same for each element for the set V_0 of clauses having variables in predicate positions (if any)
- finally, we sort the resulting set of clause numbers and hand it over to the main Prolog engine for unification and possible unfolding in case of success

Ready to run: some performance tests

Trying out the implementation

- we prototyped the design described so far as a small, slightly more than 1000 lines of “C-friendly” Java
- available at <https://github.com/ptarau/iProlog>
- for more details: recording <https://www.youtube.com/watch?v=SRyAMt8iQSw&list=PLJq3XDLIJkib2h2fObomdFRZrQeJg4UIW> of our VMSS’2016 invited tutorial
- while implemented as an interpreter, our preliminary tests indicate, very good performance

System	11 queens	perms of 11 + nrev	sudoku 4x4	metaint perms
our interpreter	5.710s	5.622s	3.500s	16.556s
Lean Prolog	3.991s	5.780s	3.270s	11.559s
Styla	13.164s	14.069s	22.196s	37.800s
SWI-Prolog	1.835s	2.620s	1.336s	4.872s
LIPS	7,278,988	7,128,483	9,261,376	6,651,000

Timings and number of logical inferences per second (LIPS) (as counted by SWI-Prolog) on 4 small Prolog programs

Summary and conclusions

- by starting from a **two line meta-interpreter**, we have captured the necessary step-by-step transformations that one needs to implement in a procedural language that mimics it
- by deriving “from scratch” a fairly efficient Prolog machine we have, hopefully, made its design more intuitive
- we have decoupled the indexing algorithm from the main execution mechanism of our Prolog machine
- we have also proposed a natural language style, human readable intermediate language that can be loaded directly by the runtime system using a minimalistic parser
- **the code and the heap representation became one and the same**
- performance of the interpreter based on our design was able to get **close enough to optimized compiled code**
- future ports of this design can help with the embedding of logic programming languages as lightweight software or hardware components