

# Formula Transformers and Combinatorial Test Generators for Propositional Intuitionistic Theorem Provers

Paul Tarau

Department of Computer Science and Engineering  
University of North Texas  
*paul.tarau@unt.edu*

**Abstract.** We develop combinatorial test generation algorithms for progressively more powerful theorem provers, covering formula languages ranging from the implicational fragment of intuitionistic logic to full intuitionistic propositional logic. Our algorithms support exhaustive and random generators for formulas of these logics. To provide known-to-be-provable formulas, via the Curry-Howard formulas-as-types correspondence, we use generators for typable lambda terms and combinator expressions. Besides generators for several classes of formulas, we design algorithms that restrict formula generation to canonical representatives among equiprovable formulas and introduce program transformations that reduce formulas to equivalent formulas of a simpler structure. The same transformations, when applied in reverse, create harder formulas that can catch soundness or incompleteness bugs. To test the effectiveness of the testing framework itself, we describe use cases for deriving lightweight theorem provers for several of these logics and for finding bugs in known theorem provers. Our Prolog implementation available at: <https://github.com/ptarau/TypesAndProofs> and a subset of formula generators and theorem provers, implemented in Python is available at: <https://github.com/ptarau/PythonProvers>.

**Keywords:** *term and formula generation algorithms, Prolog-based theorem provers, formulas-as-types, type inference and type inhabitation, combinatorial testing, finding bugs in theorem provers.*

## 1 Introduction

Theorem provers have been used in the last half-century not just to solve interesting mathematical problems but also in important practical software and hardware verification projects, ranging from nuclear reactor controllers and space-ship components to pacemakers and floating point units.

Correctness and performance of theorem provers are usually being tested using comprehensive repositories of “human-made” problems, such as the TPTP library<sup>1</sup> for classical logic or the ILTP library<sup>2</sup> for intuitionistic logic. Similarly, extensive online benchmarks help evaluate SAT, SMT and ASP solvers. Besides their chance to spot out

---

<sup>1</sup> <http://tptp.cs.miami.edu/~tptp/>

<sup>2</sup> <http://www.iltp.de/>

correctness and scalability issues, “human-made” test libraries, often derived from interesting mathematical problems, can measure the ability of theorem provers and solver engines to work well on specific problem domains.

However, as automation of testing is gaining significant traction in both software and hardware validation, it is natural to think about adopting automated testing techniques for theorem provers, which are, after all, software artifacts. Even if some of the “human-made” test sets have accumulated over the years hundreds and often thousands of problem instances, truly “adversarial” computer generated correctness and scalability tests can spot out soundness, completeness or termination issues overlooked by implementors of the intricate, heuristic-driven code of today’s theorem provers. At the same time, validation via a trusted, “gold-standard” prover, producing the same results on the same test set, can be used to propagate incrementally correctness of provers from simple formally validated versions to more sophisticated versions implementing complex heuristics.

Designing the algorithms that generate tests for theorem provers is facilitated by the regular structure of their input formulas. Such tests can be based on exhaustive small formula generators as well as random large formula generators. Besides comparison with lightweight versions of the provers, for which correctness is formally provable (possibly via proof assistants like Coq [1] or Agda [2]), known isomorphisms between formula languages and computational mechanisms like typed lambda calculi offer opportunities for transferring properties across “bridges” like the Curry-Howard *formulas-as-types* correspondence [3, 4], that ensures that the inferred type of a lambda term is a tautology in intuitionistic logic.

While the problem of finding a lambda term that has a given simple type (called the *inhabitation problem*) is PSPACE-complete [5], efficient algorithms have been known for a long time for inferring the simple type of a lambda expression, when it exists [6]. The key step in the “inner loop” of this process is *unification with occurs-check* [7], for which today’s Prolog systems offer highly efficient implementations.

The symbiosis between Automated Theorem Proving and Logic Programming has been observed in the evolution of both research fields as early as in [8]. With sound unification and backtracking efficiently implemented in today’s logic programming languages (e.g., Prolog, Curry, Picat), one can take advantage of the natural synergy that exists in these languages with features like Definite Clause Grammars (DCGs), to provide together an ideal playground for exploring combinatorial properties of typed lambda terms [9] and corresponding formula languages, essential for their applications to generation of very large terms and valid formulas. They also provide an ideal framework for transliterating sequent calculus rules into executable code.

We have built our Prolog-based open-source testing framework covering combinatorial test generators and several lightweight theorem provers for propositional intuitionistic and classical logics. Beside several of our own and 3-rd party provers, the github site<sup>3</sup> contains test generators and formula readers converting the “human-made” tests at <http://www.iltp.de> to Prolog and Python form. Part of our testing framework focusing on the implicational fragment of intuitionistic logic is described in [10] where examples of step-by step, test-driven derivations of provers are given. More re-

---

<sup>3</sup> <https://github.com/ptarau/TypesAndProofs>

cently, using these tests on the full intuitionistic propositional calculus has revealed interesting examples of Byzantine failures occurring with some of the 3-rd party provers we tested. For instance, increasing the standard 600 second timeout has revealed cases of non-termination leading to stack overflows and unexpected space complexity resulting in heap overflows, when given a very large RAM (e.g., 64GB or 96GB).

The two main generator families implemented in our testing framework are exhaustive formula generators and random formula generators. Exhaustive formula generators enumerate all formulas of a given size and thus are useful for finding minimal failure instances for a given incorrect prover. Random formulas, especially if generated as known-to-be-provable, besides pointing out soundness bugs, are also relevant as scalability tests, catching unexpected space or computation time explosion.

At the same time our testing framework contains several representation transformers that convert classes of formulas to equivalent or equiprovable<sup>4</sup> formulas.

Through a series of use cases, we exhibit provers obtained via test-driven refinements and discuss their improvements in performance and reduced space complexity.

We summarize here the main contributions of this paper:

- new combinatorial test generation algorithms for (sub-)formula languages of intuitionistic propositional logic
- restriction mechanisms limiting formula generators to one representative per class of equiprovable formulas
- transformers from disjunction-free formulas to a Nested Horn Clause form reducing space complexity from exponential to  $O(n \log(n))$
- several lightweight theorem provers obtained as a result of test-driven refinements using our formula generators
- use cases showing effectiveness of our framework in finding bugs in theorem provers

The rest of the paper is organized as follows.

Section 2 describes exhaustive generation algorithms for formulas of given (small) size, covering formulas known to be tautologies as well as arbitrary formulas. Section 3 describes algorithms generating random formulas of the same two categories. Section 4 introduces mechanisms restricting formula generators to canonical representatives of equivalence classes as well as transformations to equivalent, structurally simpler formulas. Section 5 describes test-driven refinements of provers derived from sound and complete calculi result in significant performance or space complexity improvements. Section 6 overviews our combinatorial testing framework. Section 7 shows the effectiveness of our framework in finding bugs in theorem provers. Section 8 discusses related work and section 9 concludes the paper.

## 2 Exhaustive Formula Generation Algorithms

An advantage of exhaustive testing with all formulas of a given size is that it implicitly ensures full coverage: no path is missed simply because there are no paths left unexplored.

---


<sup>4</sup> Two formulas are called equiprovable if, finding a proof for one entails the existence of a proof for the other. In particular, logically equivalent formulas are equiprovable.

**Notations and Assumptions** As we will use **Prolog** as our meta-language, our notations will be derived as much as possible from its syntax (including token types and operator definitions). Thus, variables will be denoted with uppercase letters and, as programmer's conventions final *s* letters indicate a plurality of items (e.g., when referring to the content of  $\Gamma$  contexts). We assume that the reader is familiar with basic Prolog programming, including, besides the pure Horn clause subset, well-known builtin predicates like `memberchk/2` and `select/3`, elements of higher order programming (e.g., `call/N`), and occasional use of `CUT` and `if-then-else` constructs.

Lambda terms are built using the function symbols **a/2**=application, **l/2**=lambda binder, with a logic variable as first argument and expression as second, as well as *logic variables* representing the variables of the terms.

Type expressions (also seen as implicational formulas) are built as binary trees with the function symbol **->/2** and *logic variables at their leaves*.

**Example 1** *The S combinator (left) and its type (right, with integers as leaves):*



## 2.1 The Language of Implicational Formulas

As a result of the Curry-Howard correspondence, the language of types is isomorphic with that of *the implicational fragment of intuitionistic propositional logic*, with binary trees having variables at leaf positions and the implication operator (“ $\rightarrow$ ”) at internal nodes. We will rely on the right associativity of this operator in Prolog, that matches the standard notation in type theory.

The predicate `type_skel/3` generates all binary trees with given number of internal nodes and it labels their leaves with unique logic variables. It also collects the leaf variables to a list returned as its third argument.

```
type_skel(N,T,Vs):-type_skel(T,Vs,[],N,0).

type_skel(V,[V|Vs],Vs)-->[].
type_skel((X->Y),Vs1,Vs3)-->pred,type_skel(X,Vs1,Vs2),type_skel(Y,Vs2,Vs3).
```

Type skeletons are counted by the Catalan numbers (sequence A000108 OEIS in [11]).

**Example 2** *All type skeletons for  $N=3$ .*

```
?- type_skel(3,T,_).
T = (A->B->C->D) ; T = (A-> (B->C)->D) ; T = ((A->B)->C->D) ;
T = ((A->B->C)->D) ; T = (((A->B)->C)->D) .
```

The mechanism is extended to use additional constructors from the set  $\{\sim, \&, \vee, \leftrightarrow\}$  as internal nodes of the generated trees, to cover the language of full intuitionistic propositional calculus<sup>5</sup>.

The next step toward generating the set of all type formulas is observing that logic variables define equivalence classes that correspond to partitions of the set of variables, simply by selectively unifying them.

The predicate `mpart_of/2` takes a list of distinct logic variables and generates partitions-as-equivalence-relations by unifying them “nondeterministically”. It also collects the unique variables defining the equivalence classes, as a list given by its second argument.

```
mpart_of([], []).
mpart_of([U|Xs], [U|Us]) :- mcomplement_of(U, Xs, Rs), mpart_of(Rs, Us) .
```

To implement a set-partition generator, we split a set repeatedly in subset+complement pairs with help from the predicate `mcomplement_of/2`.

```
mcomplement_of(_, [], []).
mcomplement_of(U, [X|Xs], NewZs) :-
    mcomplement_of(U, Xs, Zs),
    mplace_element(U, X, Zs, NewZs) .

mplace_element(U, U, Zs, Zs) .
mplace_element(_, X, Zs, [X|Zs]) .
```

To generate all set partitions of a set of variables of a given size, we build a list of fresh variables with Prolog’s built-in predicate `length/2` and constrain `mpart_of/2` to use them as the set to be partitioned.

```
partitions(N, Ps) :- length(Ps, N), mpart_of(Ps, _) .
```

The counts of the resulting set-partitions (Bell numbers) corresponds to the entry A000110 in [11].

**Example 3** *Set partitions of size 3 expressed as variable equalities.*

```
?- partitions(3,P).
P = [A, A, A] ; P = [A, B, A] ; P = [A, A, B] ; P = [A, B, B] ; P = [A, B, C] .
```

Hence, we can define the language of formulas in implicational intuitionistic propositional logic, among which tautologies will correspond to simple types, as being generated by the predicate `maybe_type/3`.

<sup>5</sup> at <https://github.com/ptarau/TypesAndProofs/blob/master/allFormulas.pro>

```
maybe_type(L,T,Us):-type_skel(L,T,Vs),mpart_of(Vs,Us).
```

**Example 4** *Well-formed formulas of the implicational fragment of intuitionistic propositional logic (possibly types) of size 2.*

```
?- maybe_type(2,T,_).
T = (A->A->A) ; T = (A->B->A) ; T = (A->A->B) ; T = (A->B->B) ;
T = (A->B->C) ; T = ((A->A)->A) ; T = ((A->B)->A) ; T = ((A->A)->B) ;
T = ((A->B)->B) ; T = ((A->B)->C) .
```

The sequence 2, 10, 75, 728, 8526, 115764, 1776060, 30240210 counting these formulas corresponds to the product of Catalan number of size  $n$  and Bell numbers of size  $n + 1$ , A289679 in [11].

We use these formulas to test provers that show no false negatives on known-to-be-true formulas, as described in [10]. The main issue they can reveal is if they show false positives, by succeeding on non-tautologies. This is achieved by comparing to a trusted *gold-standard* prover, e.g., one derived directly from a calculus proven sound and complete.

## 2.2 A Nested Horn Clause Tree-skeleton Generator

The generator `genHorn/3` collects leaf variables to a list, using Prolog's DCG mechanism.

```
genHorn(N,Tree,Leaves):-genHorn(Tree,N,0,Leaves,[]).

genHorn(V,N,N)-->[V].
genHorn((A:-[B|Bs]),SN1,N3)-->{succ(N1,SN1)},[A],
    genHorn(B,N1,N2),
    genHorns(Bs,N2,N3).

genHorns([],N,N)-->[].
genHorns([B|Bs],SN1,N3)-->{succ(N1,SN1)},
    genHorn(B,N1,N2),
    genHorns(Bs,N2,N3).
```

**Example 5** *generating Nested Horn Clauses*

```
?- genHorn(3,H,Vs).
H = (A:-[B, C, D]), Vs = [A, B, C, D] ;
H = (A:-[B, (C:-[D])]), Vs = [A, B, C, D] ;
H = (A:-[(B:-[C]), D]), Vs = [A, B, C, D] ;
H = (A:-[(B:-[C, D])]), Vs = [A, B, C, D] ;
H = (A:-[(B:-[(C:-[D])])]), Vs = [A, B, C, D] .
```

Interestingly, the trees corresponding to Nested Horn Clauses are enumerated by OEIS A000108, like purely implicational formulas, corresponding to Catalan numbers. Labeling of the  $N+1$  variables serving as leaves can be handled by the same partition generator we use for labeling variables of implicational formulas.

### 2.3 Typable Closed Normal Forms of Given Size

In direct relation to their computational uses, *normal forms* of simply typed lambda terms stand out. First, this is because simply typed lambda terms are strongly normalizable (i.e., their normal forms exist and are the same independently of the evaluation order). Second, because simply-typed lambda terms share their most general types (called principal types) with their normal forms. Finally, normal forms, in combination with the right size definition [12], can be described by simple CF-grammars.

Given that all formulas inhabited by a lambda terms are also inhabited by their normal forms, we can restrict ourselves to generate only typable normal forms. By generating all typable closed normal forms of a given size, we provide known-to-be-provable formulas for the implicational fragment of intuitionistic propositional logic.

The predicate `typed_nf/2`, given a size parameter `N`, iterates, on backtracking, over lambda terms in normal form `X` of size `N` and infers, on the fly, their type `T`.

```
typed_nf(N,X:T):-typed_nf(X,T,[],N,0).

pred(SX,X):-succ(X,SX).

typed_nf(l(X,E),(P->Q),Ps)-->pred,typed_nf(E,Q,[X:P|Ps]).
typed_nf(X,P,Ps)-->typed_nf_no_left_lambda(X,P,Ps).

typed_nf_no_left_lambda(X,P,[Y:Q|Ps])--> agrees(X:P,[Y:Q|Ps]).
typed_nf_no_left_lambda(a(A,B),Q,Ps)-->pred,pred,
    typed_nf_no_left_lambda(A,(P->Q),Ps),
    typed_nf(B,P,Ps).

agrees(P,Ps,N,N):-member(Q,Ps),unify_with_occurs_check(P,Q).
```

As we only need the types corresponding to provable formulas, we can omit the lambda term, resulting in a concise tautology generator for implicational intuitionistic propositional formulas:

```
impl_taut(N,T):-impl_taut(T,[],N,0).

impl_taut((P->Q),Ps)-->pred,impl_taut(Q,[P|Ps]).
impl_taut(P,Ps)-->impl_taut_no_left_lambda(P,Ps).

impl_taut_no_left_lambda(P,[Q|Ps])--> agrees(P,[Q|Ps]).
impl_taut_no_left_lambda(Q,Ps)-->pred,pred,
    impl_taut_no_left_lambda((P->Q),Ps),
    impl_taut(P,Ps).
```

**Example 6** *Implicational tautologies, after “numbering variables” as natural numbers:*

```
implTaut(N,T):-impl_taut(N,T),natvars(T).
```

```
?- implTaut(4,T).
T = (0->1->2->3->3) ;
```

```

T = (0->1->2->3->2) ;
T = (0->1->2->3->1) ;
T = (0->1->2->3->0) ;
T = (0->(0->1)->1) ;
T = ((0->1)->0->1) ;
T = (((0->0)->1)->1) .

```

**Example 7** *Counting implicational tautologies derived from typable normal forms*

```

?- countGen2(impl_taut,15,Rs) .
Rs=[1,2,3,7,17,43,129,389,1245,4274,14991,55289,210743,826136,3354509]

```

Note that the counts are not the same as OEIS A224345 which uses “natural size” of  $\lambda$ -terms, as we use here size 0 for variables, 1 for lambdas, 2 for applications.

## 2.4 Generators of Canonical forms Using Commutativity, Associativity and Idempotence of Operators

The simplest example is the generator `allSortedHorn/2` that “preemptively” ensures that bodies of Nested Horn Clauses are sorted using Prolog’s standard order, also with duplications removed, given that conjunction is idempotent<sup>6</sup>. The resulting counts match **A105633** in [11] growing with a smaller exponent than unsorted Nested Horn Clauses, which are counted by the sequence of Catalan numbers **A000108**.

```

genSortedHorn(N,Tree,Leaves):-succ(N,SN),length(Leaves,SN),
    generateSortedHorn(Tree,Leaves,[]).

generateSortedHorn(V)-->[V].
generateSortedHorn((A:-[B|Bs]))-->[A],
    generateSortedHorn(B),
    generateSortedHorns(B,Bs).

generateSortedHorns(_,[])-->[].
generateSortedHorns(B,[C|Bs])-->
    generateSortedHorn(C),
    {B@<C},
    generateSortedHorns(C,Bs).

```

This scales even more significantly in combination with a partition generator that runs first, when more frequent identical expressions, likely to get into clause bodies are eliminated:

```

allSortedHorn(N,T):-succ(N,SN),length(Vs,SN),
    natpartitions(Vs), % first, a partition generator
    genSortedHorn(N,T,Vs). % then, a sorted Horn clause generator

```

One can, by using the generator `allStrictHorn/2`, to also eliminate the “easy” Horn clauses for which the atomic head occurs in the body, to test the provers on more interesting formulas.

<sup>6</sup> <https://github.com/ptarau/TypesAndProofs/blob/master/allFormulas.pro>



Similarly, the generator `genSortedTree/3`, in combination with a partition generator, is used by `allSortedFullFormulas/2` to reduce equivalent formulas modulo associativity and commutativity of conjunction and disjunction. Other simplifications are performed at generation time, by restricting iterated negation to at most 3, as higher number of negations reduces to such equivalent formulas.

## 2.5 Generators for “uninhabitables”

With help from a theorem prover, (e.g., the predicate `hprove/1`) we can generate trees that have no inhabitants for all partitions labeling their leaves as follows:

```
unInhabitableTree(N,T):-
  genSortedHorn(N,T,Vs),
  \+ (
    natpartitions(Vs),
    hprove(T)
  ).
```

### Example 8 *Uninhabitable trees of size 5*

```
?- unInhabitableTree(5,T),nv(T).
T = (A:-[(B:-[C]), (D:-[E, F])]) ;
T = (A:-[(B:-[C, D]), (E:-[F])]) ;
T = (A:-[(B:-[C, D, E, F])]) ;
T = (A:-[(B:-[C, D, (E:-[F])])]) ;
T = (A:-[(B:-[C, (D:-[E, F])])]) ;
T = (A:-[(B:-[C, (D:-[(E:-[F])])])]) ;
T = (A:-[(B:-[(C:-[(D:-[E, F])])])]) ;
```

We can also generate leaf labelings such that no tree they are applied to, has inhabitants, as follows.

```
unInhabitableVars(N,Vs):-N>0,
  N1 is N-1,
  vpartitions(N,Vs),natvars(Vs),
  \+ (
    genSortedHorn(N1,T,Vs),
    hprove(T)
  ).
```

### Example 9 *Uninhabitable leaf labelings of size 4*

```
?- unInhabitableVars(4,Vs),nv(Vs).
Vs = [0, 1, 0, 0] ;
Vs = [0, 1, 1, 0] ;
Vs = [0, 1, 2, 0] ;
Vs = [0, 1, 0, 2] ;
Vs = [0, 1, 1, 1] ;
Vs = [0, 1, 2, 1] ;
Vs = [0, 1, 1, 2] ;
Vs = [0, 1, 2, 2] ;
Vs = [0, 1, 2, 3].
```

These are dual to similar concepts investigated for lambda terms in [13], Motzkin trees that when labeled with any de Bruijn indices result in untypable terms. Likewise, one can consider binary trees untypable with any S,K combinator labelings.

## 2.6 Some Formula Count Sequences for Small Sizes

By counting the number of solutions of our generators by increasing sizes, we obtain some interesting formula counts. We list them here together with the names of the predicates that given N as their first argument return the list of count up to N as their second argument.

- countHornTrees = A000108: Catalan numbers 1, 2, 5, 14, 42, 132, 429, 1430, 4862
- countSortedHorn = A105633: 1, 2, 4, 9, 22, 57, 154, 429, 1223, 3550, 10455, 31160, 93802, 284789
- countHorn3 = NEW: 1, 1, 2, 5, 13, 37, 109, 331, 1027, 3241, 10367, 33531, 109463
- countSortedHorn3=NEW: 1, 2, 4, 8, 20, 47, 122, 316, 845, 2284, 6264, 17337, 48424, 136196, 385548
- all implicational intuitionistic propositional calculus formulas = A289679: 1, 2, 10, 75, 728, 8526, 115764, 1776060, 30240210
- all provable implicational intuitionistic propositional calculus formulas = NEW: 0, 1, 3, 24, 201, 2201, 27406, 391379, 6215192
- countUnInhabitableTree = NEW: 1, 0, 1, 1, 4, 7, 23, 53, 163, 432, 1306
- countUnInhabitableVars = NEW: 0, 1, 1, 4, 9, 30, 122, 528, 2517, 12951, 71455

## 3 Random Formula Generation Algorithms

An advantage of random formulas of size much larger than those generated by an exhaustive enumeration at a given size, is that such formulas can be potentially harder for the provers, reveal phenomena not present at smaller sizes (e.g., unexpected space complexity), and more generally, they can test for scalability issues.

### 3.1 Random Simply-typed Terms, with Boltzmann Samplers

Once passing correctness tests, our provers need to be tested against large random terms. The mechanism is similar to the use of all-term generators.

We generate random simply-typed normal forms, using a Boltzmann sampler along the lines of that described in [14]. The code variant, adapted to our different term-size definition is at:

<https://github.com/ptarau/TypesAndProofs/blob/master/ranNormalForms.pro>. It works as follows:

```
?- ranTNF(60,XT,TypeSize).
XT = 1(1(a(a(0, 1(a(a(0, a(0, 1(...))), s(s(0))))),
      1(1(a(a(0, a(1(...), a(..., ...))), 1(0))))))
:
(A->(((A->A)- ...)->D)->D)->M)->M),
TypeSize = 34.
```

Interestingly, partly due to the fact that there's some variation in the size of the terms that Boltzmann samplers generate, and more to the fact that the distribution of types favors (as seen in the second example) the simple tautologies where an atom identical to the last one is contained in the implication chain leading to it [15, 16], if we want to use these for scalability tests, additional filtering mechanisms need to be used to statically reject type expressions that are large but easy to prove as intuitionistic tautologies.

### 3.2 Random Implicational Formulas

The generation of random implicational formulas relies on a random binary tree generator, combined with a random set partition generator.

Our code combines an implementation of Rémy's algorithm [17], along the lines of Knuth's algorithm **R** in [18] for the *generation of random binary trees* at <https://github.com/ptarau/TypesAndProofs/blob/master/RemyR.pro> with code to generate *random set partitions* at:

<https://github.com/ptarau/TypesAndProofs/blob/master/ranPartition.pro>.

We refer to [19] for a declarative implementation of Rémy's algorithm in Prolog with code adapted for this paper at:

<https://github.com/ptarau/TypesAndProofs/blob/master/RemyP.pro>.

As automatic Boltzmann sampler generation of set partitions is limited to fixed numbers of equivalence classes from which a CF- grammar can be given, we build our random set partition generator that groups variables in leaf position into equivalence classes by using an urn-algorithm [20]. Once a random binary tree of size  $N$  is generated with the  $\rightarrow/2$  constructor labeling internal nodes, the  $N + 1$  leaves of the tree are decorated with variables denoted by successive integers starting from 0. As variables sharing a name define equivalence classes on the set of variables, each choice of them corresponds to a set partition of the  $N + 1$  nodes. Thus, a set partition of the leaves  $\{0, 1, 2, 3\}$  like  $\{\{0\}, \{1, 2\}, \{3\}\}$  will correspond to the variable leaf decorations

0, 1, 1, 2

The partition generator works as follows:

```
?- ranSetPart(7,Vars).
Vars = [0, 1, 2, 1, 1, 2, 3] .
```

Note that the list of labels it generates can be directly used to decorate the random binary tree generated by Rémy's algorithm, by unifying the list of variables  $Vs$  with it.

```
?- remy(6,T,Vs).
T = (((A->B)->C->D)->E->F)->G),
Vs = [A, B, C, D, E, F, G] .
```

The combined generator, that produces in a few seconds terms of size 1000, works as follows:

```
?- time(ranImpFormula(1000,_)).
% includes tabling large Stirling numbers
% 37,245,709 inferences,7.501 CPU in
```

7.975 seconds (94% CPU, 4965628 Lips)

```
?- time(ranImpFormula(1000,_)). % fast, thanks to tabling
% 107,163 inferences,0.040 CPU in
0.044 seconds (92% CPU, 2659329 Lips)
```

Note that we use Prolog’s *tabling* (a form of automated dynamic programming) to avoid costly recomputation of the (very large) Sterling numbers in the code at: <https://github.com/ptarau/TypesAndProofs/blob/master/ranPartition.pro>.

### 3.3 Generating Random Tautologies from Typable Combinator Expressions

Boltzmann samplers for the uniform random generation of simply typed lambda terms and their normal forms are described in [14]. Of particular interest for their use as generators of random intuitionistic tautologies are the types of the terms in normal form, as every type inferred for a simply typed term can also be obtained after  $\beta$ -reduction, from its normal form. Thus, we can work on a smaller set of terms while obtaining the same set of formulas<sup>7</sup>.

One might ask why not use Hilbert-style axioms with substitutions and modus ponens to generate directly provable formulas. After all, these axioms are simple enough and exactly mimic the types of the *S* and *K* combinators:

$$\begin{aligned} K : & \quad A \rightarrow (B \rightarrow A) \\ S : & \quad (A \rightarrow (B \rightarrow C)) \rightarrow ((A \rightarrow B) \rightarrow (A \rightarrow C)) \end{aligned}$$

From them, we derive new theorems by applying substitution of formulas for variables in axioms and theorems and by applying the modus ponens inference rule:

$$MP : \quad A, A \rightarrow B \vdash B.$$

In fact, doing so would bring us back in time to the 1930’s, before Gentzen’s work on using sequent calculus for deduction [21] ! The problem is that while applying substitutions to the axioms and theorems is fairly simple (especially with Prolog’s logic variables), finding the two already known theorems needed to activate modus ponens in a growing stream of theorems is computationally prohibitive. A better (implicit) use of the *S* and *K* axioms is by designing a generator for simply typed SK-expressions. With them, one obtains the same set of the types/tautologies as those generated for simply typed lambda terms using Boltzmann samplers, given that all lambda terms are expressible as combinator formulas.

Our implementation<sup>8</sup> generates uniformly random binary trees of a given size using Rémy’s algorithm [17], along the lines of Knuth’s algorithm **R** in [18], with leaves decorated with randomly selected symbols from the set  $\{s, k\}$ . A declarative Prolog implementation of Rémy’s algorithm is described in [19]. We have also adapted its code<sup>9</sup>,

<sup>7</sup> See <https://github.com/ptarau/TypesAndProofs/blob/master/ranNormalForms.pro>

<sup>8</sup> at <https://github.com/ptarau/TypesAndProofs/blob/master/RemyR.pro>

<sup>9</sup> at <https://github.com/ptarau/TypesAndProofs/blob/master/RemyP.pro>

as a somewhat slower alternative to Knuth's algorithm. Using Knuth's algorithm **R**, the predicate `remy_sk/2` generates in a few seconds random SK-trees with 2-3 million nodes.

The predicate `ranSK/3` filters the random SK-trees of size `N` to represent typable combinator expressions, while ensuring that their types are of size at least `M`, to avoid the frequently occurring trivial types.

```
ranSK(N,M,T):-
  repeat,
    % generates a binary expression tree with S or K at its leaves
    remy_sk(N,X),
    sk_type_of(X,T), % infers the type of an SK-expression
    tsize(T,S), % computes the size of the inferred type
    S>=M,
  !,
  natvars(T). % binds type variables to natural numbers, starting from 0
```

The type inference algorithm for SK-expressions is quite simple. After stating that `s` and `k` leaves are well typed, we ensure that the types of application nodes agree (as in the modus-ponens rule), using sound unification to avoid creation of cyclical type formulas.

```
sType((A->B->C)->(A->B)->A->C).

kType((A->_B->A)).

sk_type_of(s,T):-sType(T). % S-leaf's type
sk_type_of(k,T):-kType(T). % K-leaf's type
sk_type_of((A*B),Target):- % application node
  sk_type_of(A,SourceToTarget),
  sk_type_of(B,Source),
  unify_with_occurs_check(SourceToTarget,(Source->Target)).
```

**Example 10** *Large random implicational tautologies, comparable to those generated using Boltzmann samplers, can be produced in a few seconds by inferring the types of random SK-expressions.*

```
?- ranSK(60,40,T).
T = (((((0->((1->2->3)->(1->2)->1->3)->4)->0)->0->((1->2->3)->
(1->2)->1->3)->4)->(0->((1->2->3)->(1->2)->1->3)->4)->0)->
((0->((1->2->3)->(1->2)->1->3)->4)->0)->0->((1->2->3)->(1->2)->1->3)->4)->
((1->2->3)->(1->2)->1->3)->4).
```

## 4 Formula Transformers Reducing Test Sets to Canonical Representatives of Equivalence Classes

Formula transformers serve several purposes. First, we want to perform simplifications to facilitate the work of the provers. This is achieved by converting between equivalent representations w.r.t. provability. Second, the tautologies we generate might be too easy

for the provers and we can, by applying the transformations in reverse, make them significantly harder, while still knowing their status as being provable or not.

Conversely, relying on provers known to be sound and complete, we can establish correctness of the transformers as agreement on the success of a correct prover before and after a transformation is applied.

#### 4.1 The Mints Transformation

Grigori Mints has proven, in his seminal paper studying complexity classes for intuitionistic propositional logic [22], that a formula  $f$  is equiprovable to a formula of the form  $X_f \rightarrow g$  where  $X_f$  is a conjunction of formulas of one of the forms  $p, \sim p, p \rightarrow q, (p \rightarrow q) \rightarrow r, p \rightarrow (q \rightarrow r), p \rightarrow (q \vee r), p \rightarrow \sim q, \sim q \rightarrow p$ . With introduction of new variables (like with the Tseitin transform for SAT or ASP solvers), the transformation is linear in space.

We have implemented a variant of the Mints transformation<sup>10</sup> that also eliminates negation by replacing  $\sim p$  with  $p \rightarrow \text{false}$  and expands the equivalence relation “ $\leftrightarrow$ ”. The correctness of our implementation has been tested by showing that on formulas of small sizes, a trusted prover succeeds on the same set of formulas before and after the transformation. As transforming formulas known-to-be-true results in formulas of a larger size, we have used them as scalability tests for the provers. For disjunction-free formulas, in combination with a converter to Nested Horn Clause form, the transformation has been used to generate equivalent Nested Horn Clauses of depth at most 3, a new canonical form, also useful for scalability tests for our provers.

#### 4.2 Transforming Disjunction-free Propositional Formulas to Lists of Nested Horn Clauses

The predicate `toNestedHorn/2` transforms a disjunction-free propositional formula to a Nested Horn Clause form, which is essentially the same as the language of propositional N-Prolog [23].

```
toNestedHorn(A,R):-
    expand_equiv(A,X),toHorn1(X,H),expand_horn(H,E),reduce_heads(E,R).
```

After expanding equivalences  $A \leftrightarrow B$  to conjunctions of implications it converts chained implications to Horn clauses, flattens conjunctions in their bodies and reduces formulas in head positions until all heads are atomic<sup>11</sup>. The resulting formulas can then be proven or refuted by invoking a Nested Horn Clause prover (like `ahprove/1`, to be described in section 5) on each member of the list of nested clauses. This will result in reducing worst case space complexity from exponential to  $O(n \log(n))$ .

**Example 11** *Expansion to equivalent set of Nested Horn Clauses.*

```
?-toNestedHorn(a&b&(c&d->e)<->f&g,R).
R = [(f:-[a,b,(e:-[c,d])]),(g:-[a,b,(e:-[c,d])]),(a:-[f,g]),
      (b:-[f,g]),(e:-[c,d,f,g])].
```

<sup>10</sup> <https://github.com/ptarau/TypesAndProofs/blob/master/mints.pro>

<sup>11</sup> <https://github.com/ptarau/TypesAndProofs/blob/master/toHorn.pro>

### 4.3 Transforming to the disjunction-biconditional-negation base

An alternative base for intuitionistic propositional logic is the one consisting of disjunction, biconditional and negation. Implication and conjunction can be expressed in terms of them as follows.

```

toDisjBiCond((A->B),R):-!,toDisjBiCond(A,X),toDisjBiCond(B,Y),
    R=((X v Y)<->Y).
toDisjBiCond(A & B,R):-!,toDisjBiCond(A,X),toDisjBiCond(B,Y),
    R=((X v Y)<->(X<->Y)).
toDisjBiCond(A v B,R):-!,toDisjBiCond(A,X),toDisjBiCond(B,Y),
    R=(X v Y).
toDisjBiCond(A<->B,R):-!,toDisjBiCond(A,X),toDisjBiCond(B,Y),
    R=(X<->Y).
toDisjBiCond(~A,R):-!,toDisjBiCond(A,X),
    R = (~X).
toDisjBiCond(A,A).

```

This makes formulas larger and much harder to solve, especially as biconditional “<->” is expanded to a conjunction of implications. Note that the reverse of the transformation actually works as a good simplifier for formulas passed to the provers.

## 5 Deriving Lightweight Theorem Provers for Intuitionistic Propositional Logic

Initially, like for other fields of mathematics and logic, Hilbert-style axioms were considered for intuitionistic logic. While simple and directly mapped to SKI-combinators via the Curry-Howard isomorphism, their usability for automation is very limited. In fact, their inadequacy for formalizing even “hand-written” mathematics was the main trigger of Gentzen’s work on natural deduction and sequent calculus, inspired by the need for formal reasoning in the foundation of mathematics [21].

Thus, we start with Gentzen’s own calculus for intuitionistic logic, simplified here to only cover the purely implicational fragment, given that our focus is on theorem provers working on formulas that correspond to types of simply-typed lambda terms.

### 5.1 Gentzen’s LJ Calculus, Restricted to the Implicational Fragment of Propositional Intuitionistic Logic

We assume familiarity with basic sequent calculus notation. Gentzen’s original LJ calculus [21] (with the equivalent notation of [24]) uses the following rules.

$$LJ_1 : \frac{}{A, \Gamma \vdash A}$$

$$LJ_2 : \frac{A, \Gamma \vdash B}{\Gamma \vdash A \rightarrow B}$$

$$LJ_3 : \frac{A \rightarrow B, \Gamma \vdash A \quad B, \Gamma \vdash G}{A \rightarrow B, \Gamma \vdash G}$$

As one can easily see, when trying a goal-driven implementation that uses the rules in upward direction, the unchanged premises on left side of rule  $LJ_3$  would not ensure termination as nothing prevents  $A$  and  $G$  from repeatedly trading places during the inference process.

A good starting point for developing heuristic-free, lightweight provers is to directly derive them from calculi that have been proven sound and complete.

## 5.2 The LJT/G4ip Calculus, Restricted to the Implicational Fragment

Motivated by problems related to loop avoidance in implementing Gentzen's **LJ** calculus, Roy Dyckhoff [24] introduces the following rules for his LJT calculus<sup>12</sup>.

$$LJT_1 : \frac{}{A, \Gamma \vdash A}$$

$$LJT_2 : \frac{A, \Gamma \vdash B}{\Gamma \vdash A \rightarrow B}$$

$$LJT_3 : \frac{B, A, \Gamma \vdash G}{A \rightarrow B, A, \Gamma \vdash G}$$

$$LJT_4 : \frac{D \rightarrow B, \Gamma \vdash C \rightarrow D \quad B, \Gamma \vdash G}{(C \rightarrow D) \rightarrow B, \Gamma \vdash G}$$

The rules work with the context  $\Gamma$  being either a multiset or a set.

In [10], the following literal translation of the rules  $LJT_1 \dots LJT_4$  to Prolog is given, with values in the environment  $\Gamma$  denoted by the variable  $Vs$ .

```
lprove(T):-ljt(T, []).

ljt(A,Vs):-memberchk(A,Vs),!.           % LJT_1
ljt((A->B),Vs):-!,ljt(B,[A|Vs]).         % LJT_2
ljt(G,Vs1):- %atomic(G),                 % LJT_3
    select((A->B),Vs1,Vs2),
    memberchk(A,Vs2),!,
    ljtb(G,[B|Vs2]).
ljb(G,Vs1):-                               % LJT_4
    select(((C->D)->B),Vs1,Vs2),
    ljtb((C->D),[(D->B)|Vs2]),!,
    ljtb(G,[B|Vs2]).
```

<sup>12</sup> Also called the G4ip calculus. Restricted here to the implicational fragment.



Note the use of `select/3` to extract a term from the environment (a nondeterministic step). The advantage of these rules is that they do not need loop checking to ensure termination, as one can identify a multiset ordering-based size definition that decreases after each step [24].

Next, we will show provers derived from `lprove/1` via refinements validated by our testing framework, among which ones that reduce the exponential worst case space complexity of `lprove/1` to  $O(n \log(n))$ .

### 5.3 Concentrating Nondeterminism into One Place

We start with a transformation that keeps the underlying implicational formula unchanged. It merges the work of the two `select/3` calls into a single call, observing that their respective clauses do similar things after the call to `select/3`. That avoids redoing the same iteration over candidates for reduction.

```
bprove(T):-ljb(T,[]),!.

ljb(A,Vs):-memberchk(A,Vs),!.
ljb((A->B),Vs):-!,ljb(B,[A|Vs]).
ljb(G,Vs1):-
    select((A->B),Vs1,Vs2),
    ljb_imp(A,B,Vs2),
    !,
    ljb(G,[B|Vs2]).

ljb_imp((C->D),B,Vs):-!,ljb((C->D),[(D->B)|Vs]).
ljb_imp(A,_,Vs):-atomic(A),memberchk(A,Vs).
```

### 5.4 Implicational Formulas as Nested Horn Clauses

Given the equivalence between:  $B_1 \rightarrow B_2 \dots B_n \rightarrow H$  and (in Prolog notation)  $H :- B_1, B_2, \dots, B_n$ , (where we choose  $H$  as the *atomic* formula ending a chain of implications), we can, recursively, transform an implicational formula into one built from nested clauses, as follows.

```
toHorn((A->B),(H:-Bs)):-!,toHorns((A->B),Bs,H).
toHorn(H,H).

toHorns((A->B),[HA|Bs],H):-!,toHorn(A,HA),toHorns(B,Bs,H).
toHorns(H,[],H).
```

Note also that the transformation is reversible and that lists (instead of Prolog's conjunction chains) are used to collect the elements of the body of a clause.

```
?- toHorn(((0->1->2->3->4)->(0->1->2)->0->2->3),R).
R = (3:-[(4:-[0, 1, 2, 3]), (2:-[0, 1]), 0, 2]).
```

This suggests transforming provers for implicational formulas into equivalent provers working on nested Horn clauses.

```

hprove(T0):-toHorn(T0,T),ljh(T,[],!),.

ljh(A,Vs):-memberchk(A,Vs),!.
ljh((B:-As),Vs1):-!,append(As,Vs1,Vs2),ljh(B,Vs2).
ljh(G,Vs1):-
    memberchk((G:-_),Vs1),    % atomic(G), G not in Vs1
    select((B:-As),Vs1,Vs2),  % outer select loop
    select(A,As,Bs),          % inner select loop
    ljh_imp(A,B,Vs2),          % A is an element of the body of B
    !,
    trimmed((B:-Bs),NewB),    % trim off empty bodies
    ljh(G,[NewB|Vs2]).

ljh_imp((D:-Cs),B,Vs):-!,ljh((D:-Cs),[(B:-[D])|Vs]).
ljh_imp(A,_B,Vs):-memberchk(A,Vs).

trimmed((B:-[]),R):-!,R=B.
trimmed(BBs,BBs).

```

A first improvement, ensuring quicker rejection of non-theorems is the call to `memberchk` in the 3-rd clause to ensure that our goal `G` is the head of at least one of the assumptions. Once that test is passed, the 3-rd clause works as a reducer of the assumed hypotheses. It removes from the context a clause `B:-As` and it removes from its body a formula `A`, to be passed to `ljh_imp`, with the remaining context. Should `A` be atomic, we succeed if and only if it is already in the context. Otherwise, we closely mimic rule *LJT*<sub>4</sub> by trying to prove  $A = (D:-Cs)$ , after extending the context with the assumption `B:-[D]`. Note that in both cases the context gets smaller, as `As` does not contain the `A` anymore. Moreover, should the body `Bs` end up empty, the clause is downgraded to its atomic head by the predicate `trimmed/2`. Also, by having a second `select/3` call in the third clause of `ljh`, will give `ljh_imp` more chances to succeed and commit.

Thus, besides quickly filtering out failing search branches, the nested Horn clause form of implicational logic helps bypass some intermediate steps, by focusing on the head of the Horn clause, which corresponds to the last atom in a chain of implications.

## 5.5 Propagating Back the Elimination of Non-matching Heads

We can propagate back to the implicational forms used in `bprover` the observation made on the Horn-clause form that heads (as computed below) should match at least one assumption.

```

head_of(_->B,G):-!,head_of(B,G).
head_of(G,G).

```

We can apply this to `bprove/1` as shown in the 3-rd clause of `lje`, where we can also prioritize the assumption found to have the head `G`, by placing it first in the context.

```

eprove(T):-lje(T,[],!),.

lje(A,Vs):-memberchk(A,Vs),!.
lje((A->B),Vs):-!,lje(B,[A|Vs]).

```

```
lje(G,Vs0):-
  select(T,Vs0,Vs1),head_of(T,G),!,
  select((A->B),[T|Vs1],Vs2),lje_imp(A,B,Vs2),!,
  lje(G,[B|Vs2]).

lje_imp((C->D),B,Vs):-!,lje((C->D),[(D->B)|Vs]).
lje_imp(A,_,Vs):-atomic(A),memberchk(A,Vs).
```

This brings the performance of eprove within a few percents of hprove.

## 5.6 Extracting the Proof Terms

Extracting the *proof terms* (lambda terms having the formulas we prove as types) is achieved by decorating in the code with application nodes `a/2`, lambda nodes `l/2` (with first argument a logic variable) and leaf nodes (with logic variables, same as the identically named ones in the first argument of the corresponding `l/2` nodes).

The simplicity of the predicate `eprove/1` and the fact that this is essentially the inverse of a type inference algorithm (e.g., the one in [25]) point out how the decoration mechanism works.

```
sprove(T):-ljs(X,T,[]).

ljs(X,A,Vs):-memberchk(X:A,Vs),!. % leaf variable
ljs(l(X,E),(A->B),Vs):-!,ljs(E,B,[X:A|Vs]). % lambda term
ljs(E,G,Vs1):-
  member(_:V,Vs1),head_of(V,G),!, % fail if non-tautology
  select(S:(A->B),Vs1,Vs2), % source of application
  ljs_imp(T,A,B,Vs2), % target of application
  !,
  ljs(E,G,[a(S,T):B|Vs2]). % application

ljs_imp(l(X,E),(C->D),B,Vs):-!,ljs(E,(C->D),[X:(D->B)|Vs]).
ljs_imp(E,A,_,Vs):-memberchk(E:A,Vs).
```

Thus, lambda nodes decorate *implication introductions* and application nodes decorate *modus ponens* reductions in the corresponding calculus. Note that the two clauses of `ljs_imp` provide the target node *T*. When seen from the type inference side, *T* is the type resulting from cancelling the source type *S* and the application type  $S \rightarrow T$ .

Calling `sprove/2` on the formulas corresponding to the types of the *S*, *K* and *I* combinators, we obtain:

```
?- sprove(((0->1->2)->(0->1)->0->2),X).
X = l(A, l(B, l(C, a(a(A, C), a(B, C))))) . % S
?- sprove((0->1->0),X).
X = l(A, l(B, A)) . % K
?- sprove((0->0),X).
X = l(A, A) . % I
```

### 5.7 A $O(n \log(n))$ Space Complexity Prover Implementing Hudelmaier's Calculus

In [26] a sequent calculus for intuitionistic propositional logic ensuring  $O(n \log(n))$  space complexity is introduced. We have implemented its restriction to the implicational subset as the predicate `nvprove/1`, derived in a few simple steps from `lprove/1`. The new variables are introduced by using a DCG transformation that advances a variable counter starting at 10000.

```
nvprove(T):-l_jnv(T,[],10000,_).

l_jnv(A,Vs)-->{memberchk(A,Vs)},!.
l_jnv((A->B),Vs)-->!,l_jnv(B,[A|Vs]).
l_jnv(G,Vs1)-->% atomic(G),
    {select((A->B),Vs1,Vs2)},
    l_jnv_imp(A,B,Vs2),
    !,
    l_jnv(G,[B|Vs2]).

l_jnv_imp((C->D),B,Vs)-->!,newvar(P),l_jnv(P,[C,(D->P),(P->B)|Vs]).
l_jnv_imp(A,_,Vs)-->{memberchk(A,Vs)}.

newvar(N,N,SN):-succ(N,SN).
```

Hudelmaier's algorithm achieves  $O(n \log(n))$  worst case space complexity by avoiding to duplicate the possibly large subterm  $D$  in rule  $LJT_4$ . After proving that this transformation results in a tautology if and only if the original term was provable, instead of duplicating  $D$  in the last clause of `l_jt/2`, Hudelmaier introduces a new variable  $P$ , that we implement using the DCG step `newvar/3` in the first clause of `l_jnv_imp/5`.

### 5.8 A $O(n * \log(n))$ Space Complexity Nested Horn Clause Prover

After the transformation steps shown for `hprove`, that use the fact that  $a_1 \rightarrow a_2 \dots \rightarrow a_n \rightarrow a_0$  is equivalent to  $a_0 \leftarrow a_1 \ \& \ a_2 \ \& \ \dots \ \& \ a_n$  and elimination of the interpreter wrapper by defining the predicate “ $\leftarrow$ ” directly, we activate the proof with `call(H)`, after using the transformer `toAHorn/2` to convert our tests from their implicational form to an equivalent Nested Horn Clause form.

```
ahprove(A):-toAHorn(A,H),call(H).
```

Then, the algorithm proceeds by reducing the uniformly represented Nested Horn Clauses of the form `Head <- ListOfBodyTerms`. Note also that the sequent-calculus form is not used anymore as a meta-rule, as it can be, equivalently, folded into a Nested Horn Clause form.

```
:-op(800,xfx,<-).

A<-Vs:-memberchk(A,Vs),!.
(B<-As)<-Vs1:-!,append(As,Vs1,Vs2),B<-Vs2.
```

```

G<-Vs1:- % atomic(G), G not on Vs1
  memberchk((G<-_),Vs1), % if not, we just fail
  select(B<-As,Vs1,Vs2), % outer select loop
  select(A,As,Bs), % inner select loop
  ahlj_imp(A,B,Vs2), % A element of the body of B
  !,
  atrimmed(B<-Bs,NewB), % trim empty bodies
  G<-[NewB|Vs2].

ahlj_imp(D<-Cs,B,Vs):-!, (D<-Cs)<-[B<-[D]|Vs].
ahlj_imp(A,_B,Vs):- memberchk(A,Vs).

atrimmed(B<-[],R):-!,R=B.
atrimmed(BBs,BBs).

```

A few words on the *story* that got us here. We have observed that the Nested Horn Clause prover `hprove/1` outperforms other provers (e.g., `bprove/1` by more than an order of magnitude (e.g., **121.006** seconds vs. **3221.227** seconds on terms of size **16**).

But, we have not had a convincing explanation why this is the case. The fact that a test-driven refinement step implementing Hudelmaier’s introduction of auxiliary variables brought our implication-based prover much closer in performance to the Nested Horn Clause transform, hinted towards the fact that that Hudelmaier’s optimization shares a relevant similarity with the Nested Horn Clause prover. Finally, it became clear that the duplicated formula `D` in `ahlj_imp/3`, as it occurs as the head of a clause, is atomic in the Nested Horn Clause prover and thus the space increase is bounded by the number of atoms in the original formula to be proven, without the need for introducing new variables.

## 5.9 A Lightweight Theorem Prover for Full Intuitionistic Propositional Logic

Starting from the sequent calculus for the full intuitionistic propositional logic in `LJT/G4ip` [24], to which we have also added rules for the “`<->`” relation, we obtain the following lightweight prover.

```

ljfa(T):- ljfa(T,[]).

ljfa(A,Vs):-memberchk(A,Vs),!.
ljfa(_,Vs):-memberchk(false,Vs),!.
ljfa(A<->B,Vs):-!,ljfa(B,[A|Vs]),ljfa(A,[B|Vs]).
ljfa((A->B),Vs):-!,ljfa(B,[A|Vs]).
ljfa(A & B,Vs):-!,ljfa(A,Vs),ljfa(B,Vs).
ljfa(G,Vs1):- % atomic or disj or false
  select(Red,Vs1,Vs2),
  ljfa_reduce(Red,G,Vs2,Vs3),
  !,
  ljfa(G,Vs3).
ljfa(A v B, Vs):-(!ljfa(A,Vs);ljfa(B,Vs)),!.

ljfa_reduce((A->B),_,Vs1,Vs2):-!,ljfa_imp(A,B,Vs1,Vs2).

```

```

ljfa_reduce((A & B),_,Vs,[A,B|Vs]):-!.
ljfa_reduce((A<->B),_,Vs,[(A->B),(B->A)|Vs]):-!.
ljfa_reduce((A v B),G,Vs,[B|Vs]):-ljfa(G,[A|Vs]).

ljfa_imp((C->D),B,Vs,[B|Vs]):-!,ljfa((C->D),[(D->B)|Vs]).
ljfa_imp((C & D),B,Vs,[(C->(D->B))|Vs]):-!.
ljfa_imp((C v D),B,Vs,[(C->B),(D->B)|Vs]):-!.
ljfa_imp((C<->D),B,Vs,[(C->D)->((D->C)->B)|Vs]):-!.
ljfa_imp(A,B,Vs,[B|Vs]):-memberchk(A,Vs).

```

We validate it first by testing it on the implicational subset, then against Roy Dyckhoff’s Prolog implementation<sup>13</sup>, working on formulas generated by the predicate `allSortedFullFormulas/2` up to size 12. Finally we run it on the human-made tests at <http://iltp.de> on which we get no errors, solving correctly 161 problems, with a 60 seconds timeout, compared with the 175 problems solved by Roy Dyckhoff’s heuristics-based 400 lines prover, with the same timeout<sup>14</sup>. On the other hand, the performance of `ahprove/1` is significantly better when compared with the `iLeanTap` [27], a 122 lines “lean” theorem prover<sup>15</sup> that only solves 35 problems correctly and makes 3 errors with the same 60 seconds timeout.

Among its applications, is a derivation of an embedding of Artemov and Protopopescu’s Intuitionistic Epistemic Logic [28] in Intuitionistic Propositional Logic [29], where this prover is used as an oracle for candidate definitions for epistemic operators for which theorems of the logic should hold and non-theorems fail.

## 6 The Testing Framework

Correctness can be checked by identifying false positives or false negatives. A false positive is a non-tautology that the prover proves, breaking the *soundness* property. A false negative is a tautology that the prover fails to prove, breaking the *completeness* property. While classical tautologies are easily tested (at small scale against truth tables, at medium scale with classical propositional provers and at larger scale with a SAT solver), intuitionistic provers require a more creative approach, given the absence of a finite truth-value table model.

As a first bootstrapping step, assuming that no “gold standard” prover is available, one can look at the other side of the Curry-Howard isomorphism, and rely on generators of (typable) lambda terms and generators implicational logic formulas, with results being checked against a trusted type inference algorithm.

As a next step, a trusted prover can be used as a “gold standard” to test both for false positives and negatives.

<sup>13</sup> [https://github.com/ptarau/TypesAndProofs/blob/master/third\\_party/dyckhoff\\_orig.pro](https://github.com/ptarau/TypesAndProofs/blob/master/third_party/dyckhoff_orig.pro)

<sup>14</sup> <https://github.com/ptarau/TypesAndProofs/blob/master/tester.pro>

<sup>15</sup> [https://github.com/ptarau/TypesAndProofs/blob/master/third\\_party/ileantap.pro](https://github.com/ptarau/TypesAndProofs/blob/master/third_party/ileantap.pro)

### 6.1 Finding False Negatives by Generating the Set of Simply Typed Normal Forms of a Given Size

A false negative is identified if our prover fails on a type expression known to have an inhabitant. Via the Curry-Howard isomorphism, such terms are the types inferred for lambda terms, generated by increasing sizes. In fact, this means that all implicational formulas having proofs shorter than a given number are all covered, but possibly small formulas having long proofs might not be reachable with this method that explores the search by the size of the proof rather than the size of the formula to be proven. We refer to [25] for a detailed description of efficient algorithms generating pairs of simply typed lambda terms in normal form together with their principal types. The code we use here is at: <https://github.com/ptarau/TypesAndProofs/blob/master/allTypedNFs.pro>

### 6.2 Finding False Positives by Generating All Implicational Formulas/Type Expressions of a Given Size

A false positive is identified if the prover succeeds finding an inhabitant for a type expression that does not have one.

We obtain type expressions by generating all binary trees of a given size, extracting their leaf variables and then iterating over the set of their set partitions, while unifying variables belonging to the same partition. We refer to [25] for a detailed description of the algorithms.

The code describing the all-tree and set partition generation as well as their integration as a type expression generator is at:

<https://github.com/ptarau/TypesAndProofs/blob/master/allPartitions.pro>.

We have tested the predicate `lprove/1` as well as all other provers derived from it for false negatives against simple types of terms up to size 15 (with size defined as 2 for applications, 1 for lambdas and 0 for variables) and for false positives against all type expressions up to size 7 (with size defined as the number of internal nodes).

An advantage of exhaustive testing with all formulas of a given size is that it implicitly ensures coverage: no path is missed simply because there are no paths left unexplored.

### 6.3 Testing Against a Trusted Reference Implementation

Assuming we trust an existing reference implementation (e.g., after it passes our generator-based tests), it makes sense to use it as a "gold standard". In this case, we can identify both false positives and negatives directly, as follows:

```
gold_test(N,Generator,Gold,Silver, Term, Res):-call(Generator,N,Term),
    gold_test_one(Gold,Silver,Term, Res),
    Res\=agreement.

gold_test_one(Gold,Silver,T, Res):-
    ( call(Silver,T) -> \+ call(Gold,T),
      Res = wrong_success
    ; call(Gold,T) -> % \+ Silver
```

```

    Res = wrong_failure
; Res = agreement
).
```

When specializing to a generator for all well-formed implication expressions, and using Dyckhoff's dprove/1 predicate as a gold standard, we have:

```

gold_test(N, Silver, Culprit, Unexp):-
    gold_test(N,allImpFormulas,dprove,Silver,Culprit,Unexp).
```

To test the tester, we design a prover that randomly succeeds or fails.

```

badProve(_) :- 0 == random(2).
```

We can now test lprove/1 and badprove/1 as follows:

```

?- gold_test(6,lprove,T,R).
false. % indicates that no false positive or negative is found

?- gold_test(6,badProve,T,R).
T = (0->1->0->0->0->0->0),
R = wrong_failure ;
...
?- gold_test(6,badProve,T,wrong_success).
T = (0->1->0->0->0->0->2) ;
...
```

A more interesting case is when a prover is only guilty of false positives. For instance, let's naively implement the intuition that a goal is provable w.r.t. an environment Vs if all its premises are provable, with implication introduction assuming premises and success achieved when the environment is reduced to empty.

```

badSolve(A):-badSolve(A,[]).

badSolve(A,Vs):-atomic(A),!,memberchk(A,Vs).
badSolve((A->B),Vs):-badSolve(B,[A|Vs]).
badSolve(_,Vs):-badReduce(Vs).

badReduce([]):-!.
badReduce(Vs):-select(V,Vs,NewVs),badSolve(V,NewVs),badReduce(NewVs).
```

As the following test shows, while no tautology is missed, the false positives are properly caught.

```

?- gold_test(6,badSolve,T,wrong_failure).
false.

?- gold_test(6,badSolve,T,wrong_success).
T = (0->0->0->0->0->0->1) ;
...
```



## 6.4 Testing With Large Random Terms

Testing for false positives and false negatives for random terms proceeds in a similar manner to exhaustive testing with terms of a given size.

Assuming Roy Dyckhoff’s prover as a gold standard, we can find out that our bprove/1 program can handle 20 terms of size 50 as well as the gold standard.

```
?- gold_ran_imp_test(20,100,bprove, Culprit, Unexpected).
false. % indicates no differences with the gold standard
```

In fact, the size of the random terms handled by bprove/1 makes using provers an appealing alternative to random lambda term generators in search for very large (lambda term, simple type) pairs. Interestingly, on the side of random simply typed terms, limitations come from their vanishing density, while on the other side they come from the known PSPACE-complete complexity of the proof procedures.

## 6.5 Scalability Tests

Besides the correctness and completeness test sets described so far, one might want also ensure that the performance of the derived provers scales up to larger terms. We show here a few such performance tests and refer the reader to our benchmarks at: <https://github.com/ptarau/TypesAndProofs/blob/master/bm.pro>.

Time is measured in seconds. The tables in Fig. 1 compare several provers on exhaustive “all-terms” benchmarks, derived from our correctness test.

First, we run them on the types inferred on all simply typed lambda terms of a given

Prover	Size	Positive	Mix	Total Time
lprove	13	0.979	0.261	1.24
lprove	14	4.551	5.564	10.116
lprove	15	30.014	5.568	35.583
lprove	16	3053.202	168.074	3221.277
bprove	13	0.943	0.203	1.147
bprove	14	4.461	4.294	8.755
bprove	15	32.206	4.306	36.513
bprove	16	3484.203	129.91	3614.114
dprove	13	5.299	0.798	6.098
dprove	14	23.161	13.514	36.675
dprove	15	107.264	13.645	120.909
dprove	16	1270.586	240.301	1510.887

Prover	Size	Positive	Mix	Total Time
hprove	13	1.007	0.111	1.119
hprove	14	4.413	1.818	6.231
hprove	15	20.09	1.836	21.927
<b>hprove</b>	<b>16</b>	<b>90.595</b>	<b>30.713</b>	<b>121.308</b>
eprove	13	1.07	0.132	1.203
eprove	14	4.746	2.27	7.017
eprove	15	21.562	2.248	23.81
eprove	16	97.811	43.18	140.991
sprove	13	1.757	0.173	1.931
sprove	14	8.037	2.966	11.003
sprove	15	38.266	2.941	41.208
sprove	16	188.317	54.802	243.12

**Fig. 1.** Performance of provers on exhaustive tests (faster ones in the right table)

size. Note that some of the resulting types in this case can be larger and some smaller than the sizes of their inhabitants. We place them in the column *Positive* - as they are known to be all provable.

Next, we run them on all implicational formulas of a given size, set to be about half of the former (integer part of size divided by 2), as the number of these grows much

faster. We place them in the column *Mix* as they are a mix of provable and unprovable formulas.

The predicate `hprove/1` turns out to be an overall winner, followed closely by `eprove/1` that applies to implicational forms a technique borrowed from `hprove/1` to quickly filter out failing search branches.

Testing exhaustively on small formulas, while an accurate indicator for average speed, might not favor provers using more complex heuristics or extensive preprocessing, as it is the case of Dyckhoff’s original `dprove/1`.

We conclude that early rejection via the test we have discovered in the nested Horn clause form is a clear separator between the slow provers in the left table and the fast ones in the right table, a simple and useful “mutation” worth propagating to full propositional and first order provers.

As the focus of this paper was to develop a testing methodology for propositional theorem provers, we have not applied more intricate heuristics to further improve performance or to perform better on “human-made” benchmarks or compare them on such tests with other provers, as there are no purely implicational tests among at the ILTP library [30] at <http://www.iltp.de/>. On the other hand, for our full intuitionistic propositional provers at <https://github.com/ptarau/TypesAndProofs>, as well as our Python-based ones at <https://github.com/ptarau/PythonProvers>, we have adapted the ILTP benchmarks on which we plan to report in a future paper.

## 7 A Use Case: Finding Bugs in Theorem Provers

Transformations that result in equiprovable formulas can be used to find bugs in theorem provers that escape all human-made ILTP tests, as well as our own exhaustive test on formulas of small size.

### 7.1 Catching Bugs by Hardening Implicational Formulas Known-to-be Tautologies with the Mints Transformation

We start with known tautologies in implicational fragment of intuitionistic propositional calculus obtained via the Curry-Howard correspondence as well as formulas in Full intuitionistic propositional calculus proven or disproven (by the same prover or other known to be correct prover), before applying the transformation.

Consequently, we obtain harder to prove, significantly larger formulas, for which we know that they have originated from a smaller formula with known status as provable or unprovable.

As an example, we tested the `fcube 4.1` intuitionistic propositional calculus prover [31] available at <http://www2.disco.unimib.it/fiorino/fcube.html>. It is a very nice Prolog-based prover that, in our tests, has outperformed everything else on the ILTP human-made tests. It has also passed all our tests on formulas up to size **12**.

But testing against the Mints transform finds incompleteness bugs:

```
?- small_taut_bug(4,fcube).
unexpected_failure_on
0->1->2->3->0
```

```

<=>
(nv1->0->nv2)->(nv2->1->nv3)->(nv3->2->nv4)->(nv4->3->0)->
  ((0->nv2)->nv1)->((1->nv3)->nv2)->((2->nv4)->nv3)->
  ((3->0)->nv4)->nv1

```

Fortunately, they seem to be fixed in the next version of **fCube** available from the same site.

## 7.2 Catching Bugs Using more General Intuitionistic Propositional Calculus Formulas

We can catch a bug if the “suspect” disagrees with itself on the small easy formula and its hard transformed formula.

When acting on the transformed formula, with the original seen as an oracle, a prover can be found out as unsound if it proves a non-tautology and incomplete if it fails to prove a tautology.

Thus, we can use agreement with a trusted prover running on the small formula:

```

mints_fcube(A) :- mints(A, MA), fcube(MA).

?- gold_eq_neg_test(5, mints_fcube, Culprit, Unexpected).
Culprit = ~ (0<->(1<-> ~ (1<->0))), Unexpected = wrong_failure ;
Culprit = ~ (0<->(1<-> ~ (0<->1))), Unexpected = wrong_failure ;
...

```

Note that `gold_eq_neg_test` compares behavior of a given prover against a trusted “gold standard” prover. It is more relevant for human eyes to only display the source, before the transformation is applied. In this case we find formulas containing negation and equivalence on which the prover obtained by applying of the Mints transform to the suspect fails the test, revealing the same incompleteness bug.

## 8 Related Work

The related work derived from Gentzen’s **LJ** calculus is in the hundreds if not in the thousands of papers and books. Space constraints limit our discussion to the most closely related papers, directly focusing on algorithms for implicational intuitionistic propositional logic, which, as decision procedures, ensure termination without a loop-checking mechanism.

Among them the closest are [24, 32], that we have used as starting points for deriving our provers. We have chosen to implement the **LJT** calculus directly rather than deriving our programs from Roy Dyckhoff’s Prolog code. At the same time, as in Roy Dyckhoff’s original prover, we have benefitted from the elegant, loop-avoiding rewriting step also present in Hudelmaier’s work [33, 26] and originally due to Vorobiev [34]. Similar calculi, key ideas of which made it into the Coq proof assistant’s code, are described in [35].

On the other side of the Curry-Howard isomorphism, the thesis [36], described in full detail in [37], finds and/or counts inhabitants of simple types in long normal form.

But interestingly, these algorithms have not crossed, at our best knowledge, to the other side of the Curry-Howard isomorphism, in the form of theorem provers.

Using hypothetical implications in Prolog, although all with a different semantics than Gentzen’s **LJ** calculus or its **LJT** variant, go back as early as [23, 38], followed by a series of  $\lambda$ Prolog-related publications, e.g., [39]. The similarity to the propositional subsets of N-Prolog [38] and  $\lambda$ -Prolog [39] comes from their close connection to intuitionistic logic. The hereditary Harrop formulas of [39], when restricted to their implicational subset, are more easily computable with a direct mapping to Prolog, without the need of theorem prover. While closer to an **LJ**-based calculus, the execution algorithm of [38] uses restarts on loop detection instead of ensuring termination along the lines the **LJT** calculus. In [40] backtrackable linear and intuitionistic assumptions that mimic the implication introduction rule are used, but they do not involve arbitrarily deep nested implicational formulas.

Overviews of closely related calculi, using the implicational subset of propositional intuitionistic logic are [41, 32].

For generators of random lambda terms and related functional programming constructs we refer to [42, 43]. We have shared with them the goal of achieving high-probability correctness via automated combinatorial testing. Given our specific focus on propositional provers, we have been able to use all-term and all-formula generators as well as comparison mechanisms with “gold-standard” provers. We have also taken advantage of the Curry-Howard isomorphism between types and formulas to provide an initial set of known tautologies, usable as “bootstrapping mechanism” allowing to test our provers independently from using a “gold-standard”.

Generators for closed simply-typed lambda terms, as well as their normal forms, expressed as functional programming algorithms, are given in [44], derived from combinatorial recurrences for closed terms and additional filtering for typability.

The idea of using Boltzmann samplers for generating random lambda terms was first introduced in [45]. Random lambda term generation with focus on practical uses in testing programming languages and proof assistants is covered in [43], which reports using them to find bugs in the GHC Haskell compiler.

We have used extensively Prolog as a meta-language for the study of combinatorial and computational properties of lambda terms in papers like [46, 47] covering different families of terms and properties.

The idea to use types inferred for lambda terms as formulas for testing theorem provers originates in [10]. The current paper extends this line of research to the full intuitionistic propositional logic, provides a family of algorithms for exhaustive and random tautology generators (including the combinator-based generator of random tautologies). It also describes implementation of a rich set of formula transformers, among which, the one from disjunction-free formulas to Nested Horn Clauses. This, together with the  $O(n \log(n))$ -space Nested Horn Clause prover covers the highly expressive N-Prolog subset of propositional intuitionistic logic [23].

## 9 Conclusions and Future Work

Correctness and scalability testing of theorem provers is likely to impact on their application to formal methods and proof assistants. Besides the ability to also evaluate scalability and performance of theorem provers, components of our combinatorial generation library, released as open source software, have good chances to be reused as a testing harness for theorem provers for intuitionistic, temporal, modal logic, as well as SAT, ASP or SMT solvers, with structurally similar formulas. Generators for typed lambda terms can also be reused in testing type inference algorithms in newly implemented programming languages or in wrappers adding type systems for languages like Python and Javascript. Large simply typed lambda terms can be used for performance and scalability tests for run-time systems of functional language implementations.

Future work will focus on extending our formula generation techniques to test provers for intuitionistic first order logic and some of its weaker sub-logics.

## Acknowledgement

We thank the participants to the CLA'2018 and CLA'2019 for their comments and suggestions.

## References

1. The Coq development team. *The Coq proof assistant reference manual*, 2018. Version 8.8.0.
2. Ulf Norell. Dependently Typed Programming in Agda. In *Proceedings of the 6th International Conference on Advanced Functional Programming*, AFP'08, pages 230–266, Berlin, Heidelberg, 2009. Springer-Verlag.
3. W.A. Howard. The Formulae-as-types Notion of Construction. In J.P. Seldin and J.R. Hindley, editors, *To H.B. Curry: Essays on Combinatory Logic, Lambda Calculus and Formalism*, pages 479–490. Academic Press, London, 1980.
4. Philip Wadler. Propositions as types. *Commun. ACM*, 58:75–84, 2015.
5. Richard Statman. Intuitionistic Propositional Logic is Polynomial-Space Complete. *Theor. Comput. Sci.*, 9:67–72, 1979.
6. J Roger Hindley and Jonathan P Seldin. *Lambda-calculus and combinators: an introduction*, volume 13. Cambridge University Press Cambridge, 2008.
7. J. A. Robinson. A Machine-Oriented Logic Based on the Resolution Principle. *JACM*, 12(1):23–41, 1965.
8. Larry Wos and William McCune. Automated Theorem Proving and Logic Programming: a Natural Symbiosis. *The Journal of Logic Programming*, 11(1):1 – 53, 1991.
9. Paul Tarau. A Logic Programming Playground for Lambda Terms, Combinators, Types and Tree-based Arithmetic Computations. *CoRR*, abs/1507.06944, 2015.
10. Paul Tarau. A Combinatorial Testing Framework for Intuitionistic Propositional Theorem Provers. In José Júlio Alferes and Moa Johansson, editors, *Practical Aspects of Declarative Languages - 21th International Symposium, PADL 2019, Lisbon, Portugal, January 14-15, 2019, Proceedings*, volume 11372 of *Lecture Notes in Computer Science*, pages 115–132. Springer, 2019.
11. N. J. A. Sloane. The On-Line Encyclopedia of Integer Sequences. 2018. Published electronically at <https://oeis.org/>.

12. Maciej Bendkowski, Katarzyna Grygiel, Pierre Lescanne, and Marek Zaionc. A Natural Counting of Lambda Terms. In Rusins Martins Freivalds, Gregor Engels, and Barbara Catania, editors, *SOFSEM 2016: Theory and Practice of Computer Science - 42nd International Conference on Current Trends in Theory and Practice of Computer Science, Harrachov, Czech Republic, January 23-28, 2016, Proceedings*, volume 9587 of *Lecture Notes in Computer Science*, pages 183–194. Springer, 2016.
13. Olivier Bodini and Paul Tarau. On Uniquely Closable and Uniquely Typable Skeletons of Lambda Terms. In Fabio Fioravanti and John P. Gallagher, editors, *Logic-Based Program Synthesis and Transformation, LNCS 10855*, pages 252–268. Springer International Publishing, 2018.
14. Maciej Bendkowski, Katarzyna Grygiel, and Paul Tarau. Random generation of closed simply typed  $\lambda$ -terms: A synergy between logic programming and Boltzmann samplers. *TPLP*, 18(1):97–119, 2018.
15. Antoine Genitrini, Jakub Kozik, and Marek Zaionc. Intuitionistic vs. Classical Tautologies, Quantitative Comparison. In Marino Miculan, Ivan Scagnetto, and Furio Honsell, editors, *Types for Proofs and Programs, International Conference, TYPES 2007, Cividale del Friuli, Italy, May 2-5, 2007, Revised Selected Papers*, volume 4941 of *Lecture Notes in Computer Science*, pages 100–109. Springer, 2007.
16. Zofia Kostrzycka and Marek Zaionc. Asymptotic densities in logic and type theory. *Studia Logica*, 88(3):385–403, 2008.
17. Jean-Luc Rémy. Un procédé itératif de dénombrement d’arbres binaires et son application à leur génération aléatoire. *RAIRO - Theoretical Informatics and Applications - Informatique Théorique et Applications*, 19(2):179–195, 1985.
18. Donald E. Knuth. *The Art of Computer Programming, Volume 4, Fascicle 4: Generating All Trees—History of Combinatorial Generation (Art of Computer Programming)*. Addison-Wesley Professional, 2006.
19. Paul Tarau. Declarative Algorithms for Generation, Counting and Random Sampling of Term Algebras. In *Proceedings of SAC’18, ACM Symposium on Applied Computing, PL track*, Pau, France, April 2018. ACM.
20. A.J Stam. Generation of a random partition of a finite set by an urn model. *Journal of Combinatorial Theory, Series A*, 35(2):231 – 240, 1983.
21. M. E. Szabo. The Collected Papers of Gerhard Gentzen. *Philosophy of Science*, 39(1), 1972.
22. Grigori Mints. Complexity of Subclasses of the Intuitionistic Propositional Calculus. *BIT*, 32(1):64–69, 1992.
23. Dov M. Gabbay and Uwe Reyle. N-Prolog: An extension of Prolog with hypothetical implications. I. *The Journal of Logic Programming*, 1(4):319–355, 1984.
24. Roy Dyckhoff. Contraction-free sequent calculi for intuitionistic logic. *Journal of Symbolic Logic*, 57(3):795807, 1992.
25. Paul Tarau. A Hiking Trip Through the Orders of Magnitude: Deriving Efficient Generators for Closed Simply-Typed Lambda Terms and Normal Forms. In Manuel V Hermenegildo and Pedro Lopez-Garcia, editors, *Logic-Based Program Synthesis and Transformation: 26th International Symposium, LOPSTR 2016, Edinburgh, UK, Revised Selected Papers*, pages 240–255. Springer LNCS, volume 10184, September 2017. , Best paper award.
26. J. Hudelmaier. An  $O(n \log n)$ -Space Decision Procedure for Intuitionistic Propositional Logic. *Journal of Logic and Computation*, 3(1):63–75, 1993.
27. Jens Otten. ileanTAP: An Intuitionistic Theorem Prover. In Didier Galmiche, editor, *TABLEAUX ’97*, volume 1227 of *Lecture Notes in Computer Science*, pages 307–312. Springer, 1997.
28. Sergei N. Artemov and Tudor Protopopescu. Intuitionistic Epistemic Logic. *Rew. Symb. Logic*, 9(2):266–298, 2016.

29. Paul Tarau. Modality definition synthesis for epistemic intuitionistic logic via a theorem prover. *CoRR*, abs/1907.11838, 2019.
30. Thomas Rath, Jens Otten, and Christoph Kreitz. The iltp problem library for intuitionistic logic: Release v1.1. 38:261–271, 04 2007.
31. Mauro Ferrari, Camillo Fiorentini, and Guido Fiorino. fcube: An efficient prover for intuitionistic propositional logic. In Christian G. Fermüller and Andrei Voronkov, editors, *Logic for Programming, Artificial Intelligence, and Reasoning*, pages 294–301, Berlin, Heidelberg, 2010. Springer Berlin Heidelberg.
32. Roy Dyckhoff. Intuitionistic Decision Procedures Since Gentzen. In Reinhard Kahle, Thomas Strahm, and Thomas Studer, editors, *Advances in Proof Theory*, pages 245–267, Cham, 2016. Springer International Publishing.
33. J. Hudelmaier. *A PROLOG Program for Intuitionistic Logic*. SNS-Bericht-. Universität Tübingen, 1988.
34. N. N. Vorob’ev. A new algorithm of derivability in a constructive calculus of statements. *Problems of the constructive direction in mathematics. Part 1., Trudy Mat. Inst. Steklov.*, 52:193–225, 1958.
35. Hugo Herbelin. A Lambda-Calculus Structure Isomorphic to Gentzen-Style Sequent Calculus Structure. In *Selected Papers from the 8th International Workshop on Computer Science Logic*, CSL ’94, pages 61–75, London, UK, UK, 1995. Springer-Verlag.
36. Choukri-Bey Ben-Yelles. *Type assignment in the lambda-calculus: Syntax and semantics*. PhD thesis, University College of Swansea, 1979.
37. J. Roger Hindley. *Basic Simple Type Theory*. Cambridge University Press, New York, NY, USA, 1997.
38. Dov M. Gabbay. N-PROLOG: An extension of PROLOG with hypothetical implication. II. Logical foundations, and negation as failure. *The Journal of Logic Programming*, 2(4):251–283, 1985.
39. Dale Miller and Gopalan Nadathur. *Programming with Higher-Order Logic*. Cambridge University Press, New York, NY, USA, 2012.
40. Paul Tarau, Veronica Dahl, and Andrew Fall. Backtrackable State with Linear Affine Implication and Assumption Grammars. In Joxan Jaffar and Roland H.C. Yap, editors, *Concurrency and Parallelism, Programming, Networking, and Security*, Lecture Notes in Computer Science 1179, pages 53–64, Berlin Heidelberg, December 1996. Springer.
41. Dov Gabbay and Nicola Olivetti. Goal-oriented deductions. In *Handbook of Philosophical Logic*, pages 199–285. Springer, 2002.
42. Koen Claessen and John Hughes. Quickcheck: A lightweight tool for random testing of haskell programs. *SIGPLAN Not.*, 46(4):53–64, May 2011.
43. Michal H. Palka, Koen Claessen, Alejandro Russo, and John Hughes. Testing an Optimising Compiler by Generating Random Lambda Terms. In *Proceedings of the 6th International Workshop on Automation of Software Test*, AST’11, pages 91–97, New York, NY, USA, 2011. ACM.
44. Katarzyna Grygiel and Pierre Lescanne. Counting and generating lambda terms. *J. Funct. Program.*, 23(5):594–628, 2013.
45. Katarzyna Grygiel and Pierre Lescanne. Counting and generating terms in the binary lambda calculus. *J. Funct. Program.*, 25, 2015.
46. Maciej Bendkowski, Katarzyna Grygiel, and Paul Tarau. Boltzmann Samplers for Closed Simply-Typed Lambda Terms. In Yuliya Lierler and Walid Taha, editors, *Practical Aspects of Declarative Languages - 19th International Symposium, PADL 2017, Paris, France, January 16-17, 2017, Proceedings*, volume 10137 of *Lecture Notes in Computer Science*, pages 120–135. Springer, 2017. , Best student paper award.

47. Paul Tarau. On a Uniform Representation of Combinators, Arithmetic, Lambda Terms and Types. In Elvira Albert, editor, *PPDP'15: Proceedings of the 17th international ACM SIGPLAN Symposium on Principles and Practice of Declarative Programming*, pages 244–255, New York, NY, USA, July 2015. ACM.